

# Federated Learning - Summary

Willy Fitra Hendria



# Introduction

- Introduced in a paper entitled “Communication-Efficient Learning of Deep Networks from Decentralized Data”, H. Brendan McMahan, et al
- The client’s data doesn’t need to be shared in order to learn a global model. Instead, the global model is learned by aggregating the locally-computed updates from each client devices.
- Addresses the concerns of privacy and communication costs .

# Ideal Problems for Federated Learning

Have the following properties:

- 1) Training on real-world data provides a distinct advantage
- 2) The data is privacy sensitive or large in size
- 3) Labels on the data can be inferred naturally from user interaction

# Federated Optimization

Several key properties that differ from typical distributed optimization problem:

- **Non-IID**
  - Any particular client's dataset will not be representative of the population distribution
- **Unbalanced**
  - Varying amounts of local training data
- **Massively distributed**
  - The number of clients to be much larger than the average number of examples per client
- **Limited communication**
  - Clients are frequently offline or on slow or expensive connections

# Algorithm (FederatedAveraging)

**Algorithm 1** FederatedAveraging. The  $K$  clients are indexed by  $k$ ;  $B$  is the local minibatch size,  $E$  is the number of local epochs, and  $\eta$  is the learning rate.

**Server executes:**

```
initialize  $w_0$ 
for each round  $t = 1, 2, \dots$  do
   $m \leftarrow \max(C \cdot K, 1)$ 
   $S_t \leftarrow$  (random set of  $m$  clients)
  for each client  $k \in S_t$  in parallel do
     $w_{t+1}^k \leftarrow \text{ClientUpdate}(k, w_t)$ 
   $w_{t+1} \leftarrow \sum_{k=1}^K \frac{n_k}{n} w_{t+1}^k$ 
```

**ClientUpdate**( $k, w$ ): // Run on client  $k$

```
 $\mathcal{B} \leftarrow$  (split  $\mathcal{P}_k$  into batches of size  $B$ )
for each local epoch  $i$  from 1 to  $E$  do
  for batch  $b \in \mathcal{B}$  do
     $w \leftarrow w - \eta \nabla \ell(w; b)$ 
return  $w$  to server
```

- The amount of computation is controlled by three key parameters:  $C$  (Client Fraction),  $E$  (Local Epoch), and  $B$  (Local Minibatch).
- Models are locally-trained on each 1 client:  
$$w_{t+1}^k \leftarrow \text{ClientUpdate}(k, w_t)$$
- Central server averages the resulting models  
$$w_{t+1} \leftarrow \sum_{k=1}^K \frac{n_k}{n} w_{t+1}^k$$
- With  $E=1$  and  $B=\infty$ , it corresponds to FedSGD

# Experimental Results

<b>2NN</b> <i>C</i>	IID		Non-IID	
	<i>B</i> = ∞	<i>B</i> = 10	<i>B</i> = ∞	<i>B</i> = 10
0.0	1455	316	4278	3275
0.1	1474 (1.0×)	87 (3.6×)	1796 (2.4×)	664 (4.9×)
0.2	1658 (0.9×)	77 (4.1×)	1528 (2.8×)	619 (5.3×)
0.5	— (—)	75 (4.2×)	— (—)	443 (7.4×)
1.0	— (—)	70 (4.5×)	— (—)	380 (8.6×)
<b>CNN, <i>E</i> = 5</b>				
0.0	387	50	1181	956
0.1	339 (1.1×)	18 (2.8×)	1100 (1.1×)	206 (4.6×)
0.2	337 (1.1×)	18 (2.8×)	978 (1.2×)	200 (4.8×)
0.5	164 (2.4×)	18 (2.8×)	1067 (1.1×)	261 (3.7×)
1.0	246 (1.6×)	16 (3.1×)	— (—)	97 (9.9×)

With  $B=\infty$ , there is only a small advantage in increasing  $C$ . Using smaller  $B=10$  shows a significant improvement in using  $C \geq 0.1$ , especially in the non-IID case

# Experimental Results

MNIST CNN, 99% ACCURACY					
CNN	$E$	$B$	$u$	IID	Non-IID
FEDSGD	1	$\infty$	1	626	483
FEDAVG	5	$\infty$	5	179 (3.5 $\times$ )	1000 (0.5 $\times$ )
FEDAVG	1	50	12	65 (9.6 $\times$ )	600 (0.8 $\times$ )
FEDAVG	20	$\infty$	20	234 (2.7 $\times$ )	672 (0.7 $\times$ )
FEDAVG	1	10	60	34 (18.4 $\times$ )	350 (1.4 $\times$ )
FEDAVG	5	50	60	29 (21.6 $\times$ )	334 (1.4 $\times$ )
FEDAVG	20	50	240	32 (19.6 $\times$ )	426 (1.1 $\times$ )
FEDAVG	5	10	300	20 (31.3 $\times$ )	229 (2.1 $\times$ )
FEDAVG	20	10	1200	18 (34.8 $\times$ )	173 (2.8 $\times$ )


  

SHAKESPEARE LSTM, 54% ACCURACY					
LSTM	$E$	$B$	$u$	IID	Non-IID
FEDSGD	1	$\infty$	1.0	2488	3906
FEDAVG	1	50	1.5	1635 (1.5 $\times$ )	549 (7.1 $\times$ )
FEDAVG	5	$\infty$	5.0	613 (4.1 $\times$ )	597 (6.5 $\times$ )
FEDAVG	1	10	7.4	460 (5.4 $\times$ )	164 (23.8 $\times$ )
FEDAVG	5	50	7.4	401 (6.2 $\times$ )	152 (25.7 $\times$ )
FEDAVG	5	10	37.1	192 (13.0 $\times$ )	41 (95.3 $\times$ )

MNIST 2NN	$E$	$B$	$u$	IID	Non-IID
FEDSGD	1	$\infty$	1	1468	1817
FEDAVG	10	$\infty$	10	156 (9.4 $\times$ )	1100 (1.7 $\times$ )
FEDAVG	1	50	12	144 (10.2 $\times$ )	1183 (1.5 $\times$ )
FEDAVG	20	$\infty$	20	92 (16.0 $\times$ )	957 (1.9 $\times$ )
FEDAVG	1	10	60	92 (16.0 $\times$ )	831 (2.2 $\times$ )
FEDAVG	10	50	120	45 (32.6 $\times$ )	881 (2.1 $\times$ )
FEDAVG	20	50	240	39 (37.6 $\times$ )	835 (2.2 $\times$ )
FEDAVG	10	10	600	34 (43.2 $\times$ )	497 (3.7 $\times$ )
FEDAVG	20	10	1200	32 (45.9 $\times$ )	738 (2.5 $\times$ )

With  $C=0.1$ , adding more local updates per round (increase  $E$  & decrease  $B$ ) can produce a dramatic decrease in communication costs

# Experimental Results




For very large  $E$ , FedAvg can plateau or diverge.

It may be useful to decay the amount of local computation per round (moving to smaller  $E$  or larger  $B$ )




# Experimental Results



Acc.	80%		82%		85%	
SGD	18000	(—)	31000	(—)	99000	(—)
FedSGD	3750	(4.8×	6600	(4.7×	N/A	(—)
FedAVG	280	(64.3×	630	(49.2×	2000	(49.5×

On the CIFAR-10 dataset, FedAVG has less number of communication rounds compares to FedSGD and baseline SGD.

# Experimental Results



On a large-scale LSTM experiment, FedSGD with  $\eta = 18.0$  required 820 rounds to reach 10.5%, while FedAvg with  $\eta=9.0$  reached an accuracy of 10.5% in only 35 rounds, which is 23X fewer than fedSGD.